

Deep Learning: Project presentation

Deep Reinforcement Learning with JAVAgario

DEBES Baptiste

NELISSEN Louis

WEYDERS Pierre-Francois



August 2020


The Game: Agar.io

- Popular Massively Multiplayer browser game
- Player is a cell that eats pellets
- Constantly shrinks, but grows when eating pellets
- Controlled with mouse




The Game: JAVAgar.io

- Coded in JAVA
- Greyscale image
- Ability to launch our own servers
- Simplified rules
 - One agent
 - One box (small)
 - 20 pellets
- Objective: get all the pellets as fast as possible without dying



The Networks



The Theory: Q-Learning

- Branch of reinforcement learning
- Model-free
- Value based learning
 - $Q \rightarrow (\text{state}, \text{action})$
- Q-learning aims to approximate the optimal Q^*
- Deep Q-learning: use NN as approximators

The Theory: Policy Gradient

- Branch of reinforcement learning
- Model-free
- Seek an approximator $\pi_{\theta}(a|s)$ of the optimal policy $\pi^*(a|s)$

The Theory: The RL family

Value based

DDQN

- No policy (implicit)
- Value function

Policy based

- Policy
- No value function

Actor critic

DDPG

A2C

- Policy
- Value function

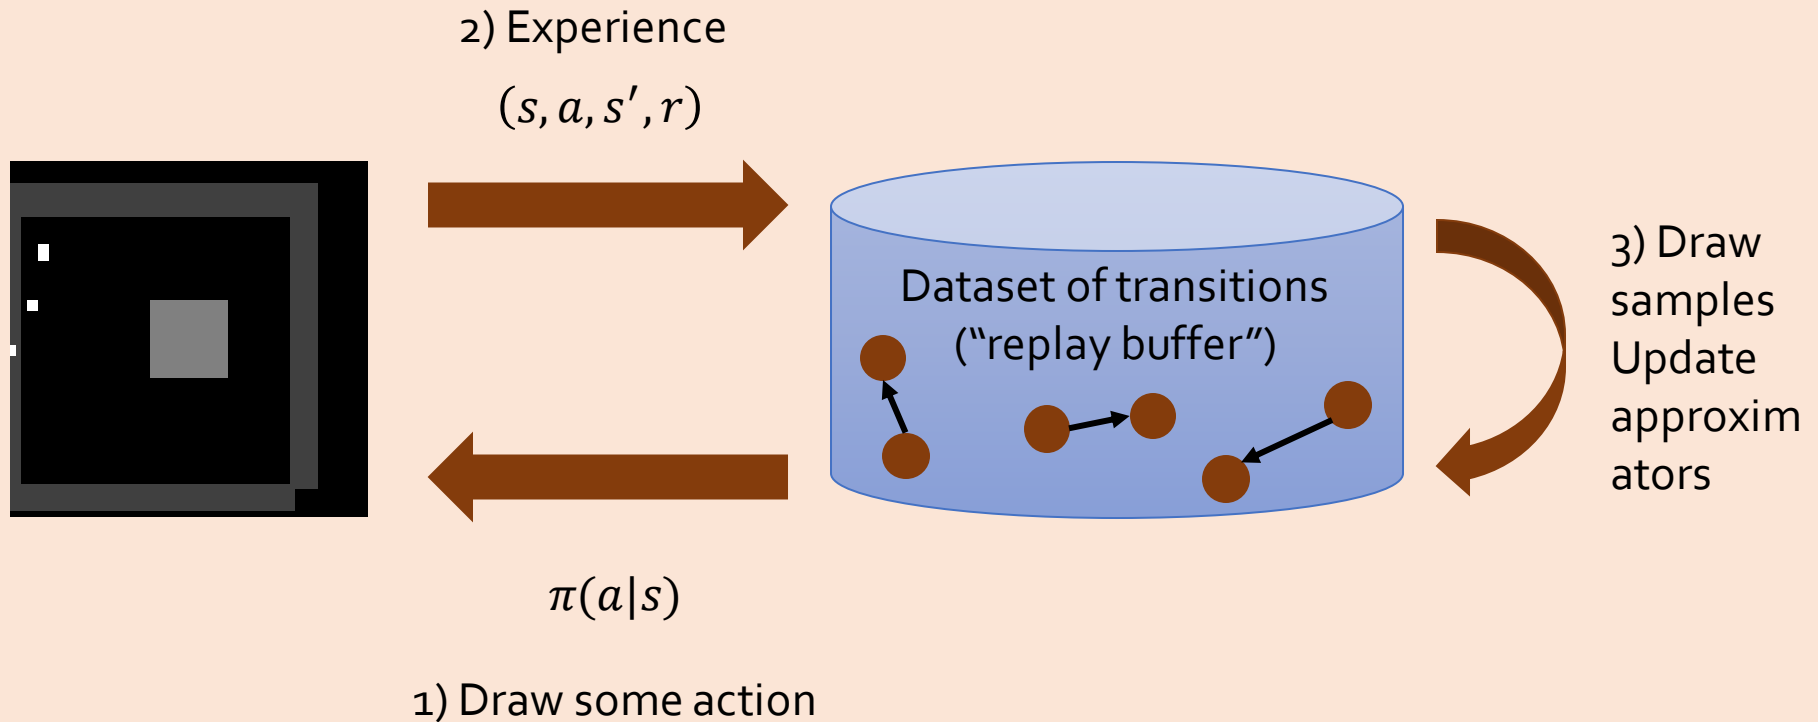
Model based

- Transition and reward model

Model free

- No transition and no reward model (implicit)

The Theory: The general algorithm



The Theory: The losses


A2C: Advantage Actor Critic

- $L^{PG}(\theta) = E \left[-\log(\pi_{\theta}(a_n|s_n)A(s_n, a_n)) \right]$

DDPG: Double DPG

- $L^{DDPG}(\theta) = -E \left[Q_{\tau}(s, a) \mid_{s=s_n, a_n=\pi_{\theta}(s_n)} \right]$

Caveat in
the report



DDQN: Double Deep Q Network


- $L^{TD}(\theta) = \frac{1}{2} \left[Q_{\theta}(s, a) - r - \gamma \max_{\{a'\}} Q_{\bar{\theta}}(s', a') \right]^2$

The Theory: Other tricks

- Batch normalization:
 - Observed in the DDPG paper
- Target networks:
 - Stabilizes training
 - Soft update equation

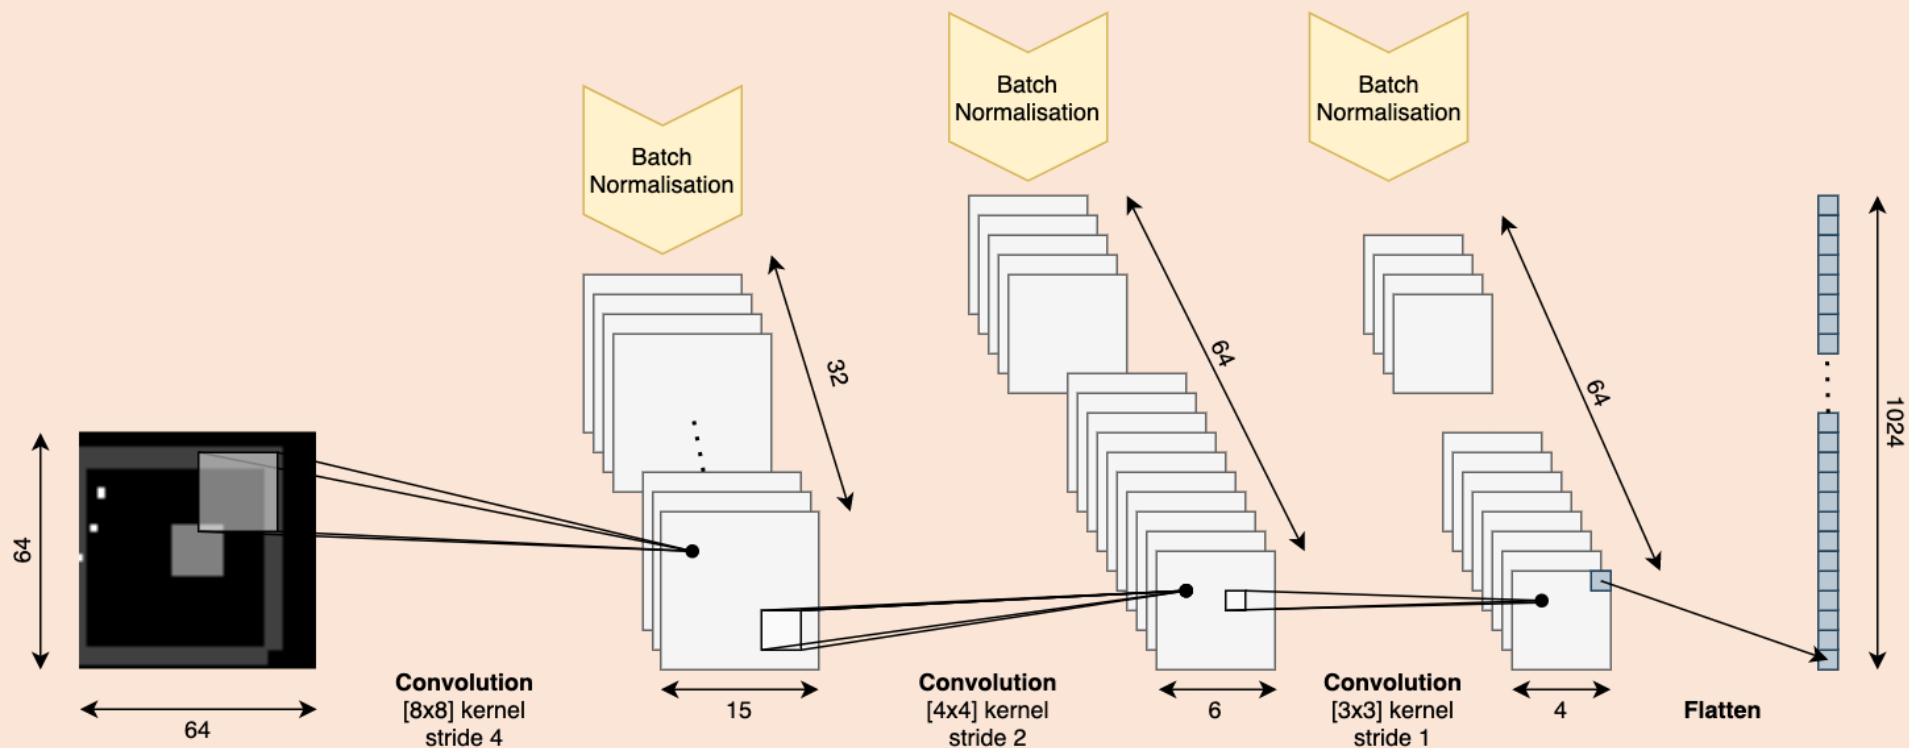
$$\theta' \leftarrow \tau * \theta + (1 - \tau) * \theta' \text{ with } \tau \in [0,1]$$

The Methods: General Network

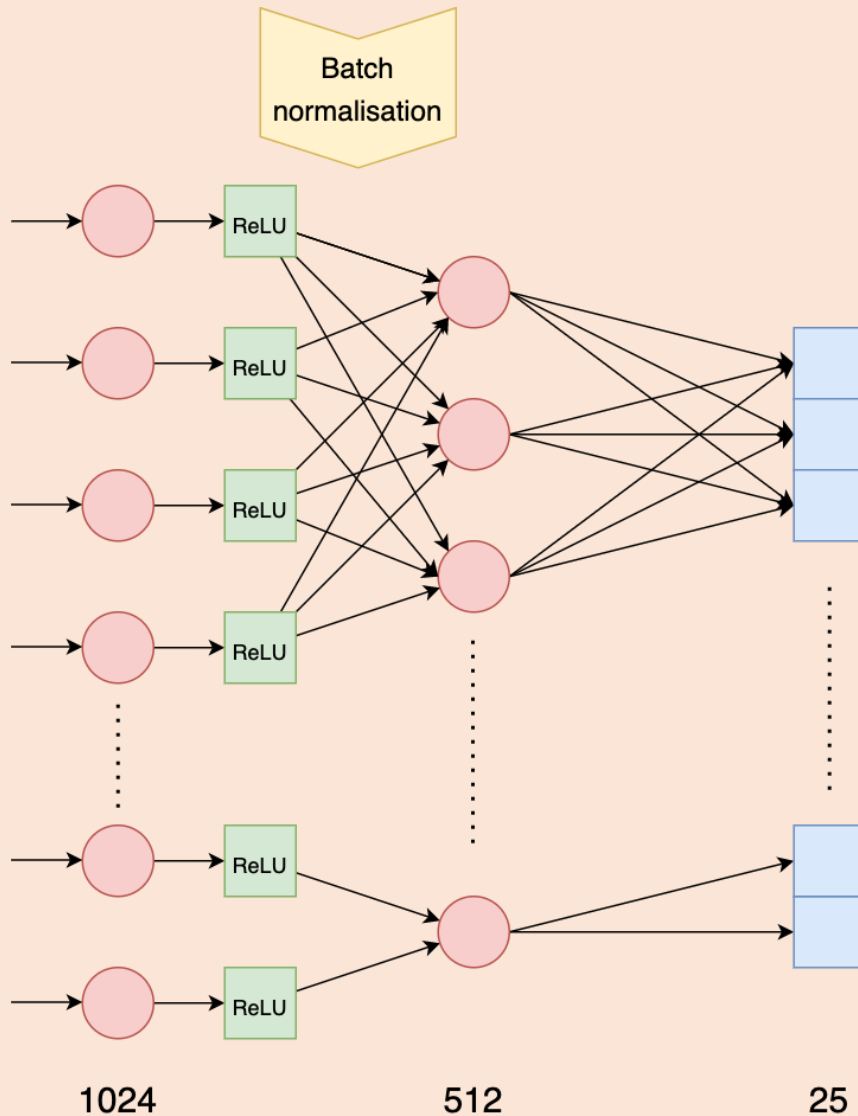


The Methods: Convolutional Network

- 3 Convolution layers
- Batch normalisation



The Methods: Multi Layer Perceptron




- 2 dense layers
- Batch normalisation


Different versions for the different algorithms

The Methods: A2C

Actor:




Critic:




Parameters	Value
Learning rate	0.0001
Tau	0.001
Discount factor	0.99

The Methods: DDPG

Actor:




Critic:



Parameters	Value
Learning rate critic	0.0005
Learning rate actor	0.00001
Tau	0.001
Discount factor	0.99

The Methods: DDQN



Parameters	Value
Learning rate	0.0001
Tau	0.001
Discount factor	0.99

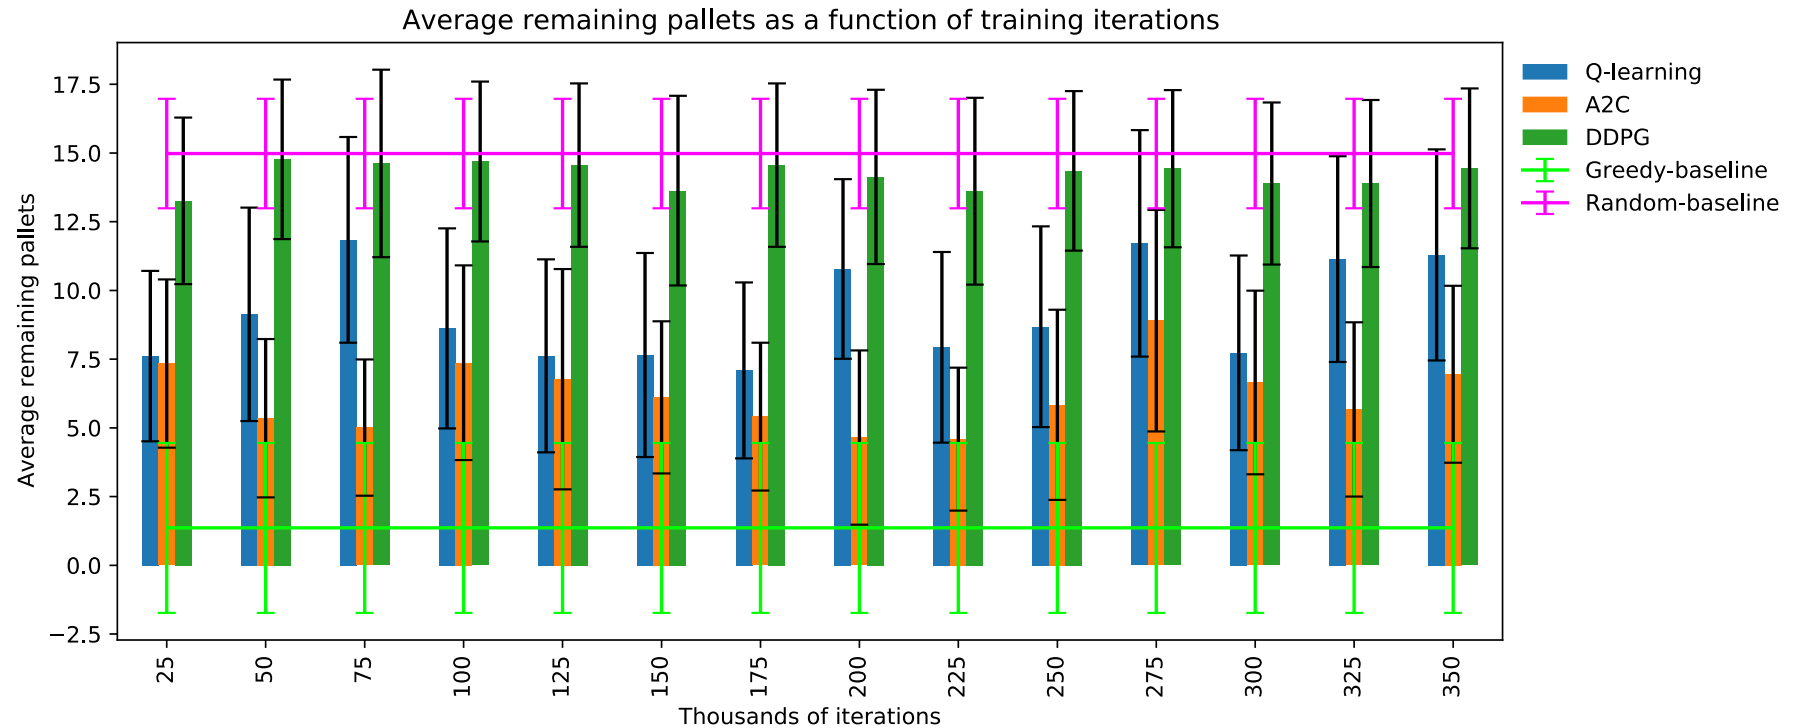
The Methods: Training

- Random environments
 - Random pellet position
 - Random init position
- Constants:
 - Size of map
 - Number of pellets (20)
- Images of 64×64
- Episodes of 2500 steps
- Total of 350 000 frames
- Replay buffer of size 65 000
- Adam optimizer


Hardware:

- *CPU*: I7-6700k 4GHZ (4 cores)
- *GPU*: NVIDIA GTX 1070 (training was ~ 30 % faster)


The Results




The Results




The Results: example trajectories



Greedy agent



A2C after 3000 steps



A2C after 4000 steps

The Results: Discussion

- Lack of training
- Simplicity of RL Algorithms
- Complexity of the input
- Size of the replay buffer

Thank you for your attention